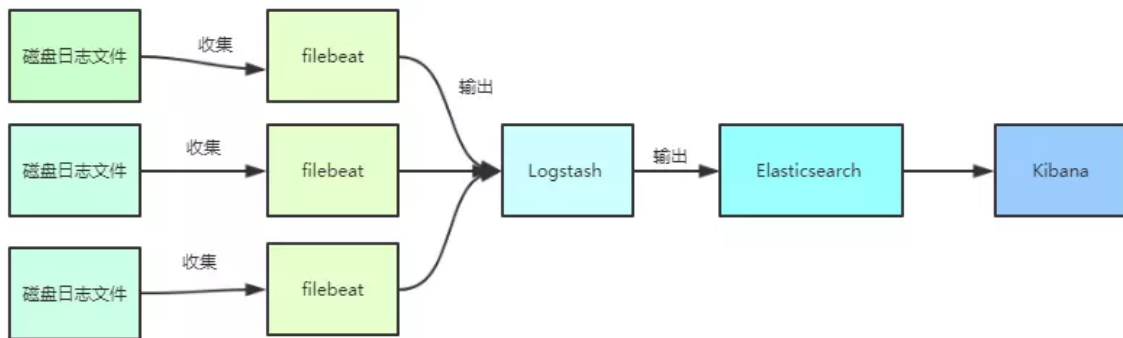


ELK是集中式、独立的、搜集管理各个服务和服务器上的日志信息，可以通过可视化的方式查看日志，帮助开发人员和运维人员快速处理和分析线上的问题。



- 1、Filebeat负责收集应用写到磁盘上的日志，并将日志发送给logstash
- 2、logstash处理来自filebeat的日志，并将处理后的日志保存elasticsearch索引库。
- 3、elasticsearch存储来自logstash的日志。
- 4、kibana从elasticsearch搜索日志，并展示到页面。

下面介绍一下整个日志管理系统的搭建过程。

环境准备：centos7, java8, node-v9.8.0 (kibana依赖nodejs)

准备好Filebeat, logstash, elasticsearch, kibana, 可以去官网下载，这里用的版本是6.2.4

## 1、安装elasticsearch。

出于安全考虑，elasticsearch默认不允许以root账号运行。

创建用户：

```
useradd suzhe
```

设置密码：

```
passwd suzhe
```

切换用户：

```
su - suzhe
```

上传安装包,并解压

```
[suzhe@VM_108_39_centos ~]$ tar -zxvf elasticsearch-6.2.4.tar.gz
```

目录重命名：

```
mv elasticsearch-6.2.4/ elasticsearch
```

进入查看目录

```
[suzhe@VM_108_39_centos elasticsearch]$ ll
total 240
drwxr-xr-x  2 suzhe suzhe  4096 Nov  8 09:59 bin
drwxr-xr-x  2 suzhe suzhe  4096 Dec 21 10:31 config
drwxrwxr-x  3 suzhe suzhe  4096 Nov  8 10:07 data
drwxr-xr-x  2 suzhe suzhe  4096 Apr 13  2018 lib
-rw-r--r--  1 suzhe suzhe 11358 Apr 13  2018 LICENSE.txt
drwxr-xr-x  2 suzhe suzhe  4096 Dec 20 05:38 logs
drwxr-xr-x 16 suzhe suzhe  4096 Apr 13  2018 modules
-rw-r--r--  1 suzhe suzhe 191887 Apr 13  2018 NOTICE.txt
drwxr-xr-x  3 suzhe suzhe  4096 Nov  8 10:40 plugins
-rw-r--r--  1 suzhe suzhe  9268 Apr 13  2018 README.textile
[suzhe@VM_108_39_centos elasticsearch]$
```

修改配置

cd config 进入配置目录。

```
[suzhe@VM_108_39_centos elasticsearch]$ cd config
[suzhe@VM_108_39_centos config]$ ll
total 16
-rw-rw----  1 suzhe suzhe 2878 Nov 29 20:39 elasticsearch.yml
-rw-rw----  1 suzhe suzhe 2771 Nov  8 10:03 jvm.options
-rw-rw----  1 suzhe suzhe 5091 Apr 13  2018 log4j2.properties
[suzhe@VM_108_39_centos config]$
```

jvm.options (Elasticsearch基于Lucene的, 而Lucene底层是java实现, 因此可以调整jvm参数)

```
-Xms1g
-Xmx1g
```

修改elasticsearch.yml: vim elasticsearch.yml

- 修改数据和日志目录:

```
path.data: /home/suzhe/elasticsearch/data # 数据目录位置
path.logs: /home/suzhe/elasticsearch/logs # 日志目录位置
```

我们把data和logs目录修改指向了elasticsearch的安装目录。但是这两个目录并不存在, 因此我们需要创建出来。

进入elasticsearch的根目录, 然后创建:

```
mkdir data
mkdir logs
```

- 修改绑定的ip:

```
network.host: 0.0.0.0 # 绑定到0.0.0.0, 允许任何ip来访问
```

默认只允许本机访问，修改为0.0.0.0后则可以远程访问

运行：

进入elasticsearch/bin目录可以看到elasticsearch 可执行文件

```
[suzhe@VM_108_39_centos elasticsearch]$ cd bin
[suzhe@VM_108_39_centos bin]$ ll
total 256
-rwxr-xr-x 1 suzhe suzhe 1557 Apr 13 2018 elasticsearch
-rw-r--r-- 1 suzhe suzhe 1431 Apr 13 2018 elasticsearch.bat
-rwxr-xr-x 1 suzhe suzhe 2238 Apr 13 2018 elasticsearch-env
-rw-r--r-- 1 suzhe suzhe 1713 Apr 13 2018 elasticsearch-env.bat
-rwxr-xr-x 1 suzhe suzhe 239 Apr 13 2018 elasticsearch-keystore
-rw-r--r-- 1 suzhe suzhe 329 Apr 13 2018 elasticsearch-keystore.bat
-rwxr-xr-x 1 suzhe suzhe 229 Apr 13 2018 elasticsearch-plugin
-rw-r--r-- 1 suzhe suzhe 319 Apr 13 2018 elasticsearch-plugin.bat
-rw-r--r-- 1 suzhe suzhe 8018 Apr 13 2018 elasticsearch-service.bat
-rw-r--r-- 1 suzhe suzhe 104448 Apr 13 2018 elasticsearch-service-mgr.exe
-rw-r--r-- 1 suzhe suzhe 103936 Apr 13 2018 elasticsearch-service-x64.exe
-rwxr-xr-x 1 suzhe suzhe 242 Apr 13 2018 elasticsearch-translog
-rw-r--r-- 1 suzhe suzhe 332 Apr 13 2018 elasticsearch-translog.bat
[suzhe@VM_108_39_centos bin]$
```

执行命令启动:

```
./elasticsearch -d
```

访问: <http://node:9200/> 可以看到如下的json信息。

```
{
  name: "cQBepWL",
  cluster_name: "elasticsearch",
  cluster_uuid: "DJkZxqH2Tpu5-uD01U4iig",
  version: {
    number: "6.2.4",
    build_hash: "ccec39f",
    build_date: "2018-04-12T20:37:28.497551Z",
    build_snapshot: false,
    lucene_version: "7.2.1",
    minimum_wire_compatibility_version: "5.6.0",
    minimum_index_compatibility_version: "5.0.0"
  },
  tagline: "You Know, for Search"
}
```

## 2、安装kibana

### 1、解压

```
[root@VM_108_39_centos software]# tar -zxvf kibana-6.2.4-linux-x86_64
```

### 2、修改配置

配置elasticsearch的地址

```
# Kibana is served by a back end server. This setting specifies the port to use.
server.port: 5601
```

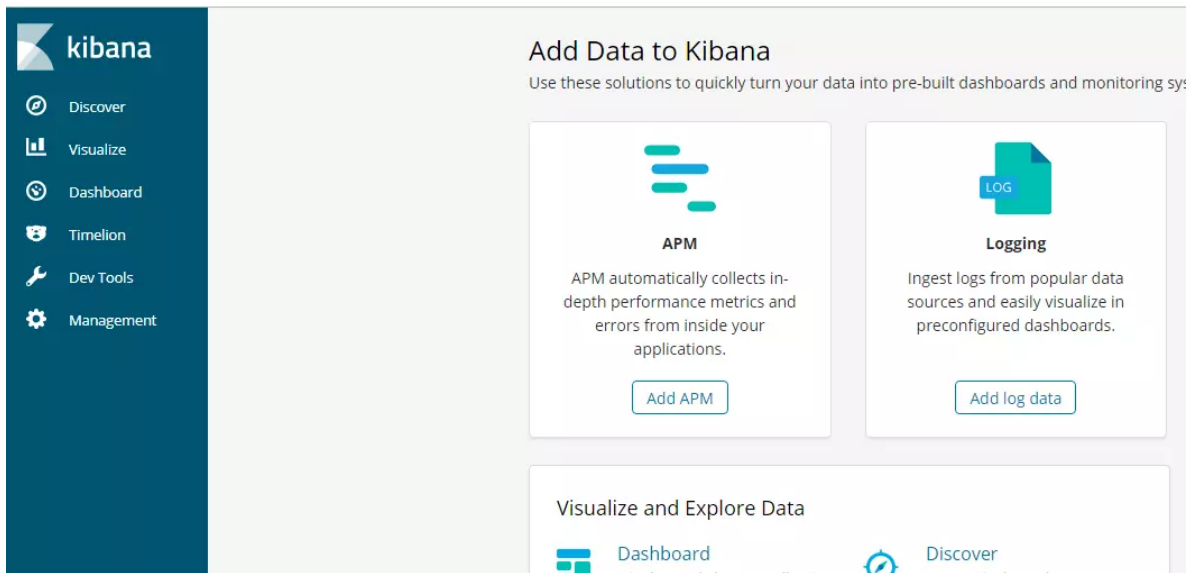
```
# To allow connections from remote users, set this parameter to a non-loopback
address.
server.host: "0.0.0.0"
```

```
# The URL of the Elasticsearch instance to use for all your queries.
elasticsearch.url: "http://localhost:9200"
```

后台运行 Kibana:

```
[root@VM_108_39_centos software]# cd kibana-6.2.4-linux-x86_64/bin/
[root@VM_108_39_centos bin]# nohup ./kibana &
```

访问: <http://node:5601>



### 3、安装Logstash

解压logstash

```
[root@VM_108_39_centos software]# tar -zxvf logstash-6.2.4.tar.gz
```

创建配置文件

```
[root@VM_108_39_centos software]# vim logstash-6.2.4/config/filebeat-first.conf
```

内容如下

```
input {
  stdin {}
  beats {
    port => 5044
  }
}
```

```

}
output {
  elasticsearch {
    hosts => ["192.168.1.56:9200"]
  }
  stdout {
    codec => rubydebug
  }
}
}

```

port是接受filebeat的端口，192.168.1.56:9200为elasticsearch的服务地址，这儿要替换成你的地址。

启动logstash

```

[root@VM_108_39_centos software]# cd logstash-6.2.4/bin/
[root@VM_108_39_centos bin]# nohup ./logstash -f ../config/filebeat-first.conf &

```

查看进程

```

[root@VM_108_39_centos bin]# ps -ef|grep logstash
root      7278   3659  99 11:30 pts/1    00:01:01 /usr/local/software/jdk1.8.0_101/bin/java -Xms1g -Xmx1g -XX:+UseParNewGC -XX:+UseConcMarkSweepGC -XX:CMSInitiatingOccupancyFraction=70 -XX:+HeapDumpOnOutOfMemoryError -Djava.secd
file:/dev/urandom -cp /usr/local/software/logstash-6.2.4/logstash-core/lib/jars/commons-compiler-3.0.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/google-java-fre
jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/quava-19.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/jackson-annotations-2.9.1.jar:/usr/local/sof
stash-6.2.4/logstash-core/lib/jars/jackson-core-2.9.1.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/jackson-databind-2.9.1.jar:/usr/local/software/logstash-6.2.4/l
core/lib/jars/jackson-dataformat-cbor-2.9.1.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/janino-3.0.8.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars
mplete-9.1.13.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/log4j-api-2.9.1.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/log4j-core-2.9.1.jar:/u
software/logstash-6.2.4/logstash-core/lib/jars/log4j-slf4j-impl-2.9.1.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/logstash-core.jar:/usr/local/software/logstash-
stash-core/lib/jars/org.eclipse.core.commands-3.6.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.core.contenttype-3.4.100.jar:/usr/local/software/Logs
4/logstash-core/lib/jars/org.eclipse.core.expressions-3.4.300.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.core.filesystem-1.3.100.jar:/usr/local/soft
stash-6.2.4/logstash-core/lib/jars/org.eclipse.core.jobs-3.5.100.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.core.resources-3.7.100.jar:/usr/local/sof
stash-6.2.4/logstash-core/lib/jars/org.eclipse.core.runtime-3.7.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.equinox.app-1.3.100.jar:/usr/local/soft
stash-6.2.4/logstash-core/lib/jars/org.eclipse.equinox.common-3.6.0.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.equinox.preferences-3.4.1.jar:/usr/loc
re/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.equinox.registry-3.5.101.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.jdt.core-3.10.0.jar:/usr/lo
are/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.osgi-3.7.1.jar:/usr/local/software/logstash-6.2.4/logstash-core/lib/jars/org.eclipse.text-3.5.101.jar:/usr/local/software/Log
4/logstash-core/lib/jars/slf4j-api-1.7.25.jar:org.logstash.logstash -f ../config/filebeat-first.conf

```

查看端口是否被监听：

```

[root@VM_108_39_centos bin]# netstat -lntp|grep 5044
tcp        0      0 0.0.0.0:5044          0.0.0.0:*            LISTEN
7278/java
[root@VM_108_39_centos bin]#

```

## 4、安装Filebeat

解压

```

[root@VM_108_39_centos software]# tar -zxvf filebeat-6.2.4-linux-x86_64.tar.gz

```

编辑配置文件

```

[root@VM_108_39_centos software]# vim filebeat-6.2.4-linux-x86_64/filebeat.yml

#===== Filebeat prospectors =====

filebeat.prospectors:

- type: log
  enabled: true
  paths:

```

```
- /data/logs/admin.log
- /data/logs/mobile.log
fields:
  log_source: node1
  logtype: applog
document_type: applog
multiline.pattern: '^\[ '
multiline.negate: true
multiline.match: after
exclude_lines: ['DEBUG']
```

```
#----- Logstash output -----
output.logstash:
  # The Logstash hosts
  hosts: ["192.168.1.56:5044"]
```

filebeat.yml 配置的主要有两个部分，一个是日志收集，一个是日志输出的配置。

配置解释：

type: log 读取日志文件的每一行（默认） enabled: true 该配置是否生效,如果改为false,将不收集该配置的日志 paths: 要抓取日志的全路径 fields: 自定义属性,可以定义多个,继续往下排就行  
multiline.pattern: 正则表达式 multiline.negate: true 或 false; 默认是false, 匹配pattern的行合并到上一行; true, 不匹配pattern的行合并到上一行 multiline.match: after 或 before, 合并到上一行的末尾或开头

exclude\_lines: ['DEBUG'] 该属性配置不收集DEBUG级别的日志,如果配置多行 这个配置也要放在多行的后面

192.168.1.56:5044 为输出到Logstash的地址和端口。

启动filebeat

```
[root@VM_108_39_centos filebeat-6.2.4-linux-x86_64]# nohup ./filebeat -e -c
filebeat.yml &
```

## 5、验证

### 1、输入日志文件

进入/data/logs目录输入日志。

```
[root@VM_108_39_centos logs]# echo "删除用户" >> admin.log
[root@VM_108_39_centos logs]# echo "提现成功 " >> mobile.log
[root@VM_108_39_centos logs]# echo "注册成功 " >> mobile.log
[root@VM_108_39_centos logs]# echo "I love you,admin" >> admin.log
```

### 2、创建index

Discover
Visualize
Dashboard
Timelion
Dev Tools
Management

Management / Kibana

Index Patterns
Saved Objects
Advanced Settings

Warning
No default index pattern. You must select or create one to continue.

### Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch inc

Step 1 of 2: Define index pattern

Index pattern

logstash-2018.12.21

You can use a \* as a wildcard in your index pattern. You can't use empty spaces or the characters \, /, ?, ", <, >, |.

✓ **Success!** Your index pattern matches 1 index.

logstash-2018.12.21

Rows per page: 10

点击next step

Index Patterns
Saved Objects
Advanced Settings

Warning
No default index pattern. You must select or create one to continue.

### Create index pattern

Kibana uses index patterns to retrieve data from Elasticsearch indices for things like visualizations.

Step 2 of 2: Configure settings

You've defined **logstash-2018.12.21** as your index pattern. Now you can specify some settings before we create it.

Time Filter field name
Refresh

@timestamp

The Time Filter will use this field to filter your data by time.

You can choose not to have a time field, but you will not be able to narrow down your data by a time range.

> Show advanced options

点击 create index pattern 创建索引

点击 discover可查看日志

Discover
Visualize
Dashboard
Timelion
Dev Tools
Management

7 hits

logstash-2018.12.21

Selected Fields
Available Fields

@timestamp
@version
\_id
\_index
\_score
\_type
beat.hostname
beat.name
beat.version
fields.log\_source
fields.logtype
host
message

December 21st 2018, 12:47:20.469 - December 21st 2018, 13:02:20.469 —

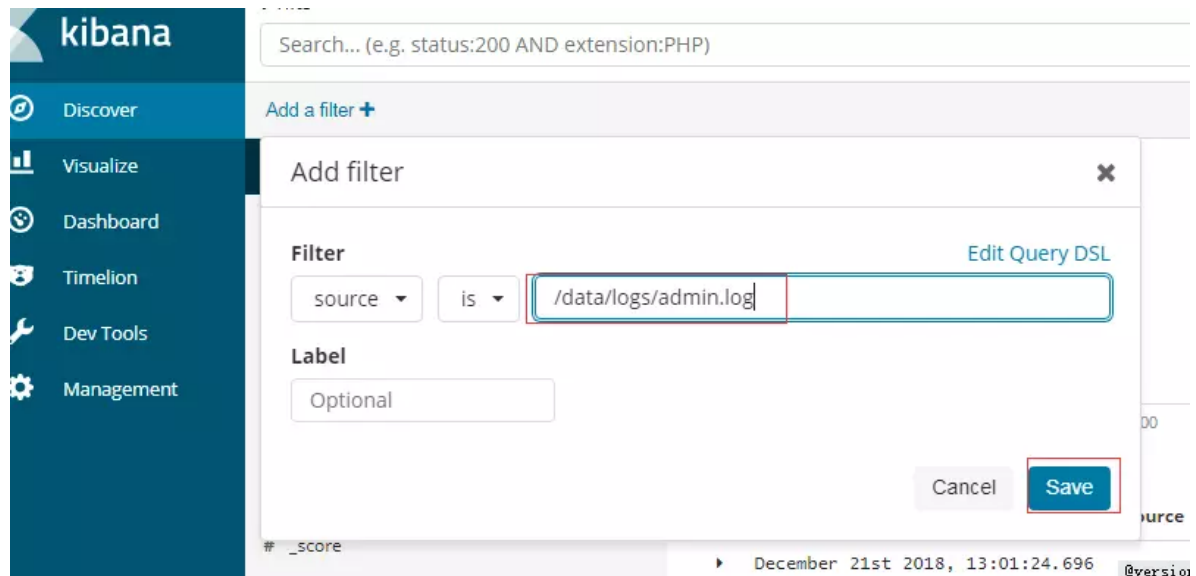
Time
\_source

December 21st 2018, 13:01:24.696
@version: 1
source: /data/logs/admin.log
message: I love you,admin
beat.name: VM
timestamp: December 21st 2018, 13:01:24.696
tags: beats\_input\_codec\_plain\_appl
offset: 43
prospector.type: log
\_id: \_4ojz2cB4E-r01uv05s7
\_type: doc
\_index: log

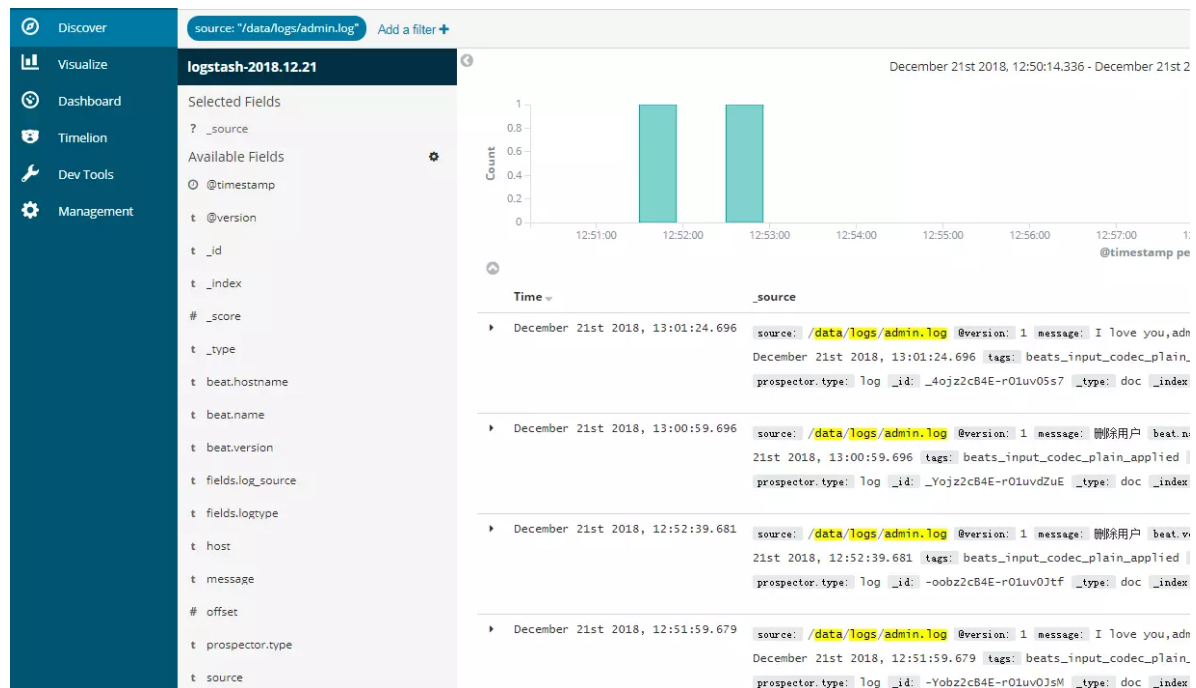
December 21st 2018, 13:01:14.696
@version: 1
source: /data/logs/mobile.log
message: 注册成功
beat.version: 6.2.4
timestamp: December 21st 2018, 13:01:14.696
tags: beats\_input\_codec\_plain\_appl
prospector.type: log
\_id: \_oojz2cB4E-r01uvrJsm
\_type: doc
\_index: logstash-2018.

December 21st 2018, 13:00:59.696
@version: 1
source: /data/logs/admin.log
message: 删除用户
beat.name: VM\_108\_39-c
timestamp: December 21st 2018, 13:00:59.696
tags: beats\_input\_codec\_plain\_appl
prospector.type: log
\_id: \_Yojz2cB4E-r01uvdZuE
\_type: doc
\_index: logstash-2018.

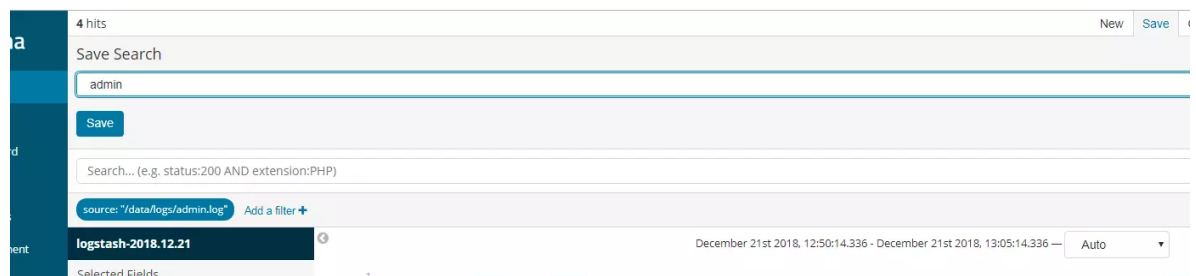
根据source过滤，只查看admin的日志




点击保存，可以看到只有admin的日志。



保存该过滤条件




下次再查看日志可以直接点击open--》点击admin进行查看





同理mobile也可以配置好过滤，下次查看日志的时候直接根据情况选择admin或者mobile。



ok，到这里，整个日志管理系统就搭建好了。

## 6、思考

以上对于一般的中小型公司，上面的系统基本上可以满足需求，只是需要做到高可用（logstash集群，elasticsearch集群，kibana高可用），这部分比较简单。但是对于高并发场景，可能会产生大量的日志，大量的数据涌入Logstash集群以及elasticsearch集群，可能系统会遇到流量上的瓶颈。

如何解决这个问题？

加入数据缓冲层，日志采集客户端采集来的数据,转存到kafka+ookeeper集群中,做一个消息队列,让数据有一定的缓冲。如图所示：

